

BAREFOOT NETWORKS

Labs Guide

11/13/2015

Environment Introduction

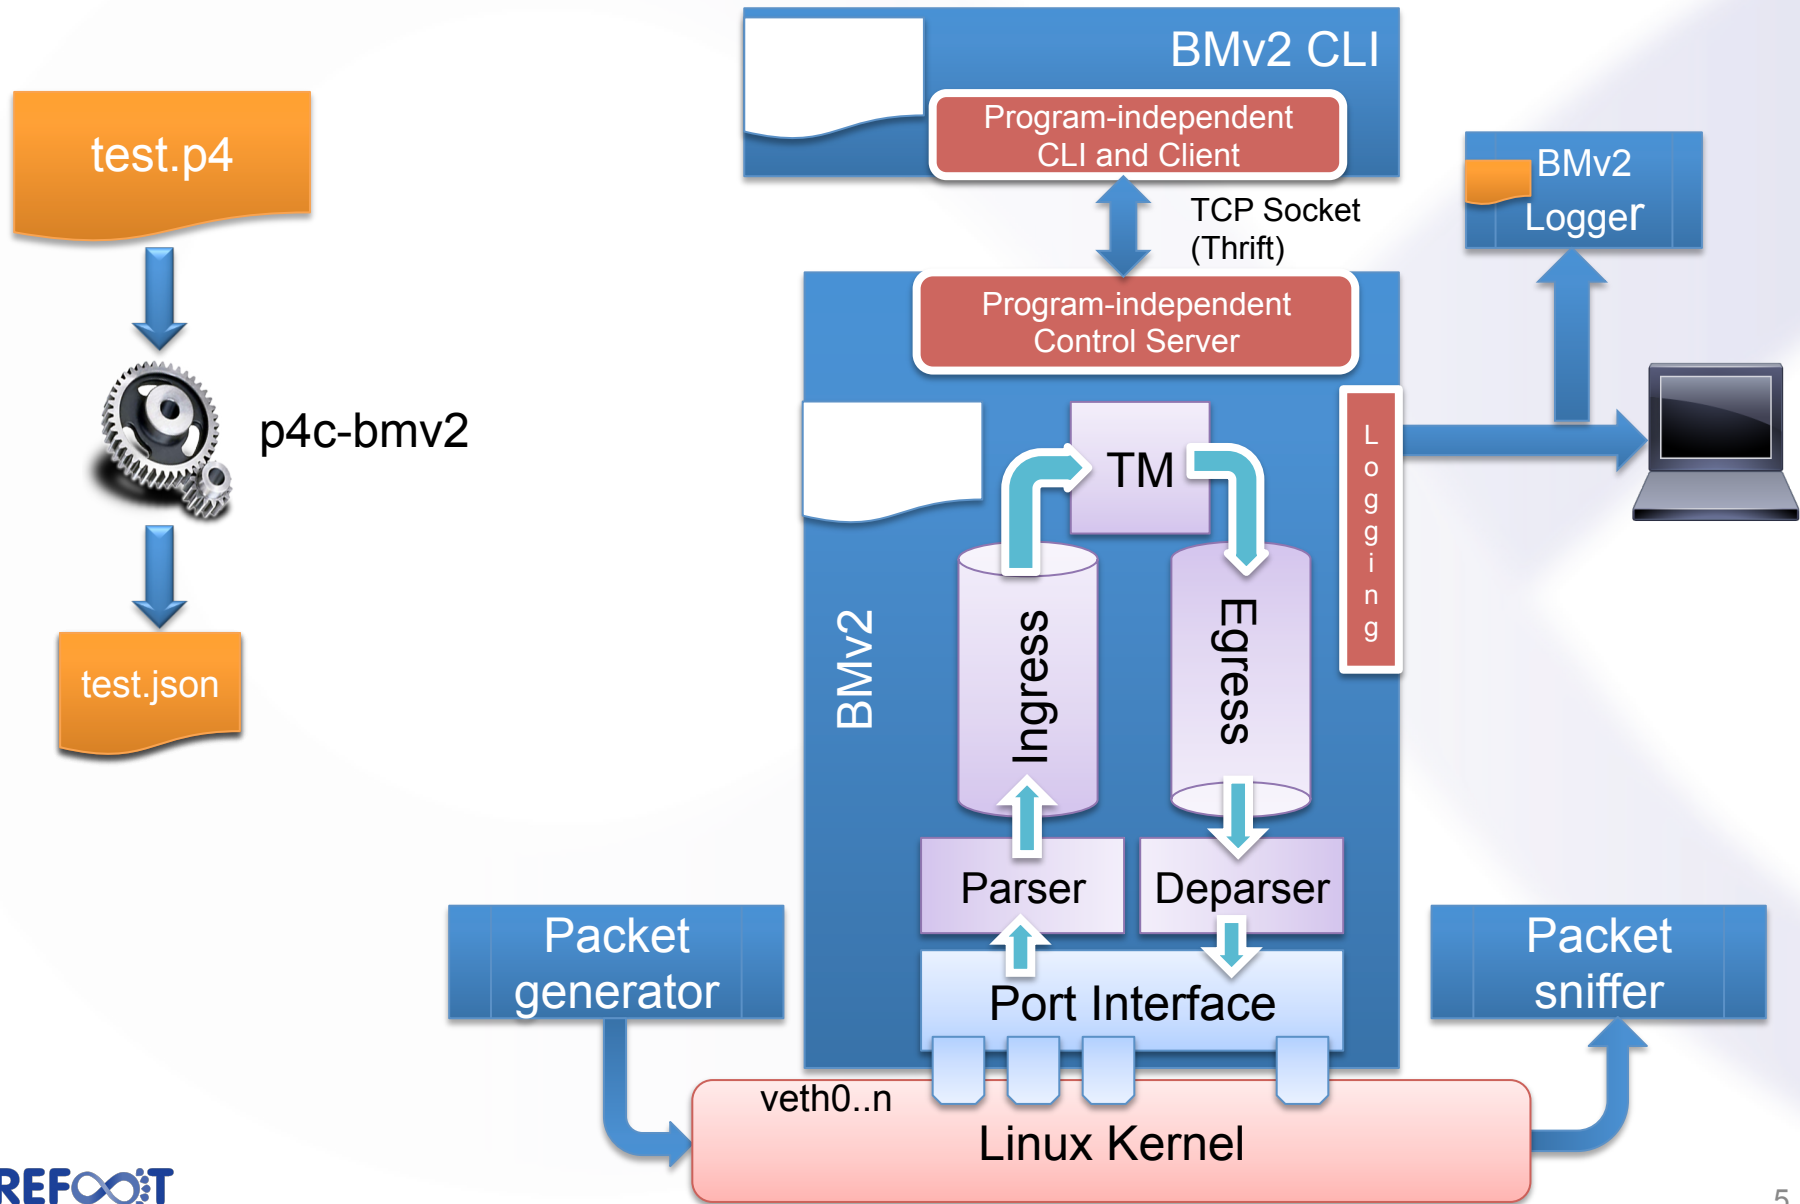
Lab VM

- **Operating System: Ubuntu 14.04**
- **Shipped as OVA (Open Virtual Appliance)**
- **Needs to be imported into Virtual Box**
 - File → Import Appliance...
- **Download link**
 - https://www.dropbox.com/s/wyjdi2bb5fv7mue/ubuntu14_04-p4bootcamp.ova?dl=0
- **User Name: ubuntu**
- **Password: ubuntu**
- **Do not forget to install VirtualBox additions!**
 - Devices → Insert Guest Additions CD Image...

Basic Workflow

- How Everything “clicks” together

Basic Workflow



Step 1: P4 Program Compilation

test.p4

```
$ p4c-bmv2 --json test.json test.p4
```




p4c-bmv2




test.json

Step 2: Preparing veth Interfaces



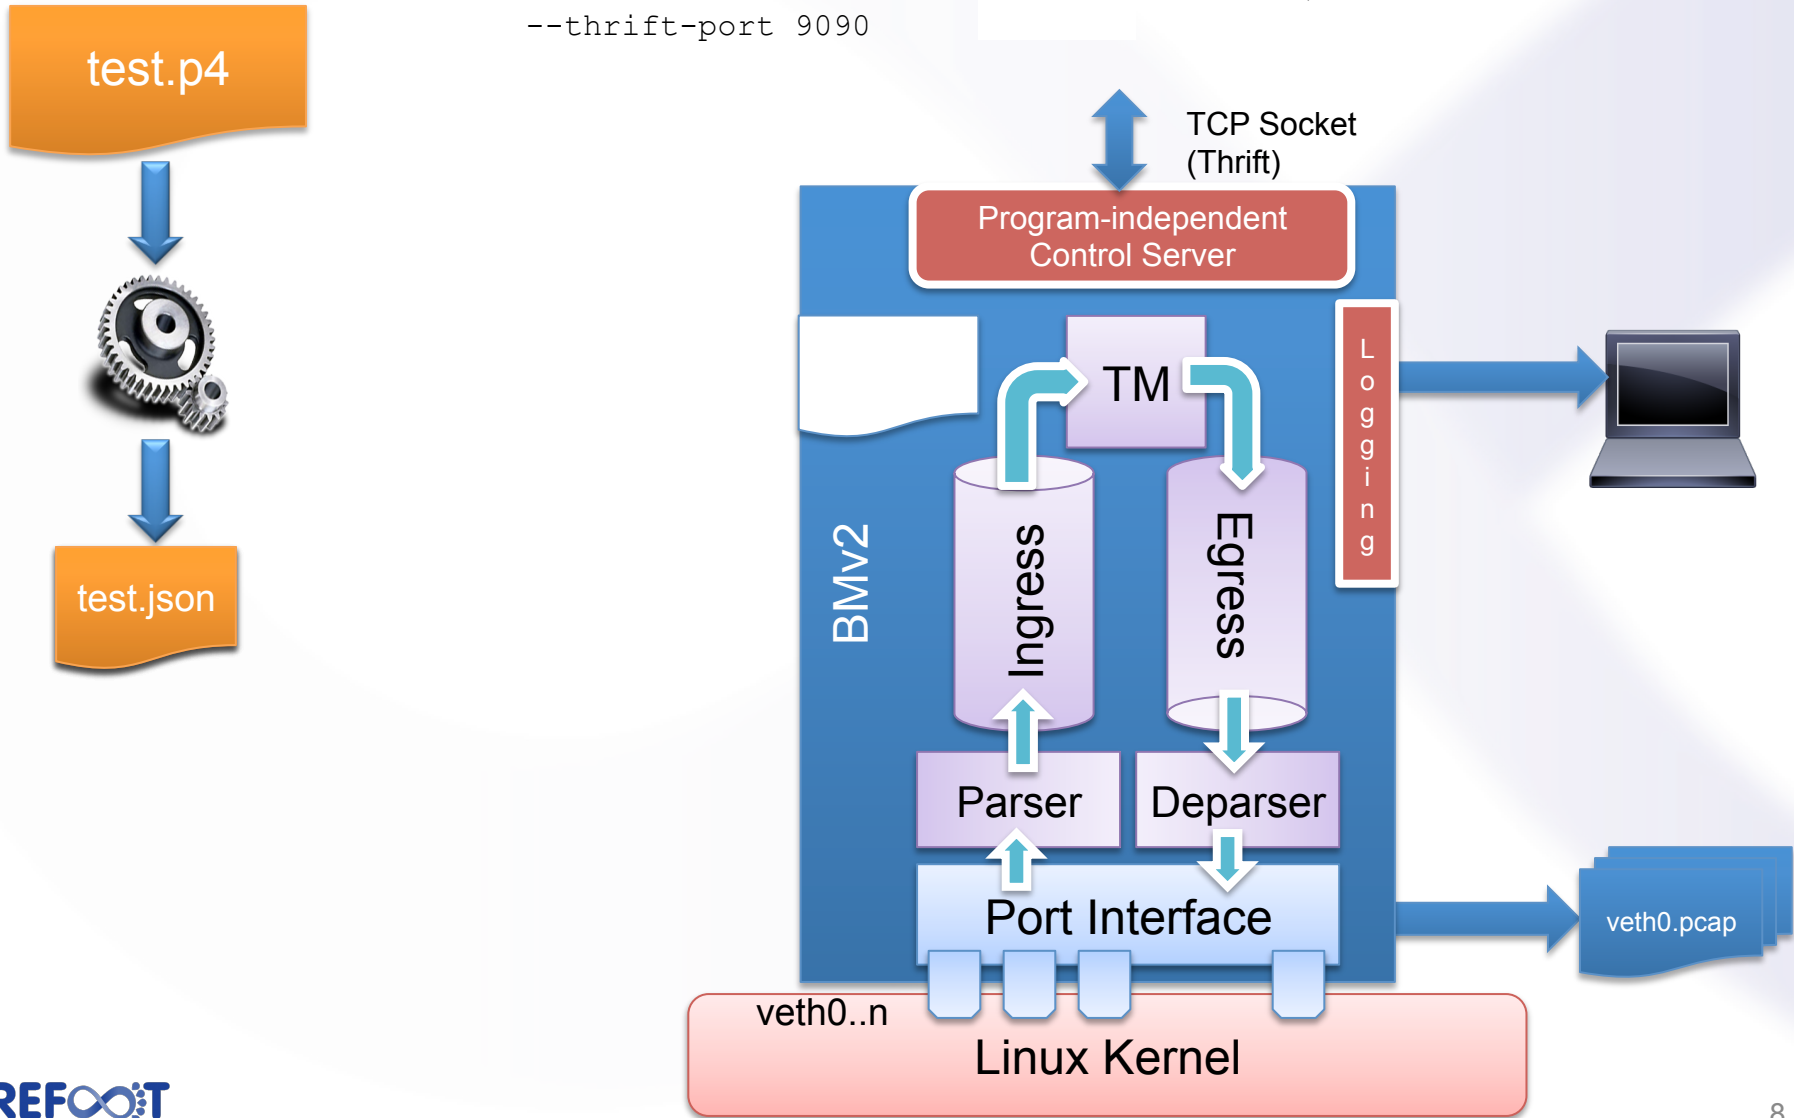
```
$ sudo ~/tutorial/examples/veth_setup.sh
```

```
# ip link add name veth0 type veth peer name veth1
# for iface in "veth0 veth1"; do
    ip link set dev ${iface} up
    sysctl net.ipv6.conf.${iface}.disable_ipv6=1
    TOE_OPTIONS="rx tx sg tso ufo gso gro lro rxvlan txvlan rxhash"
    for TOE_OPTION in $TOE_OPTIONS; do
        /sbin/ethtool --offload $intf "$TOE_OPTION"
    done
done
```



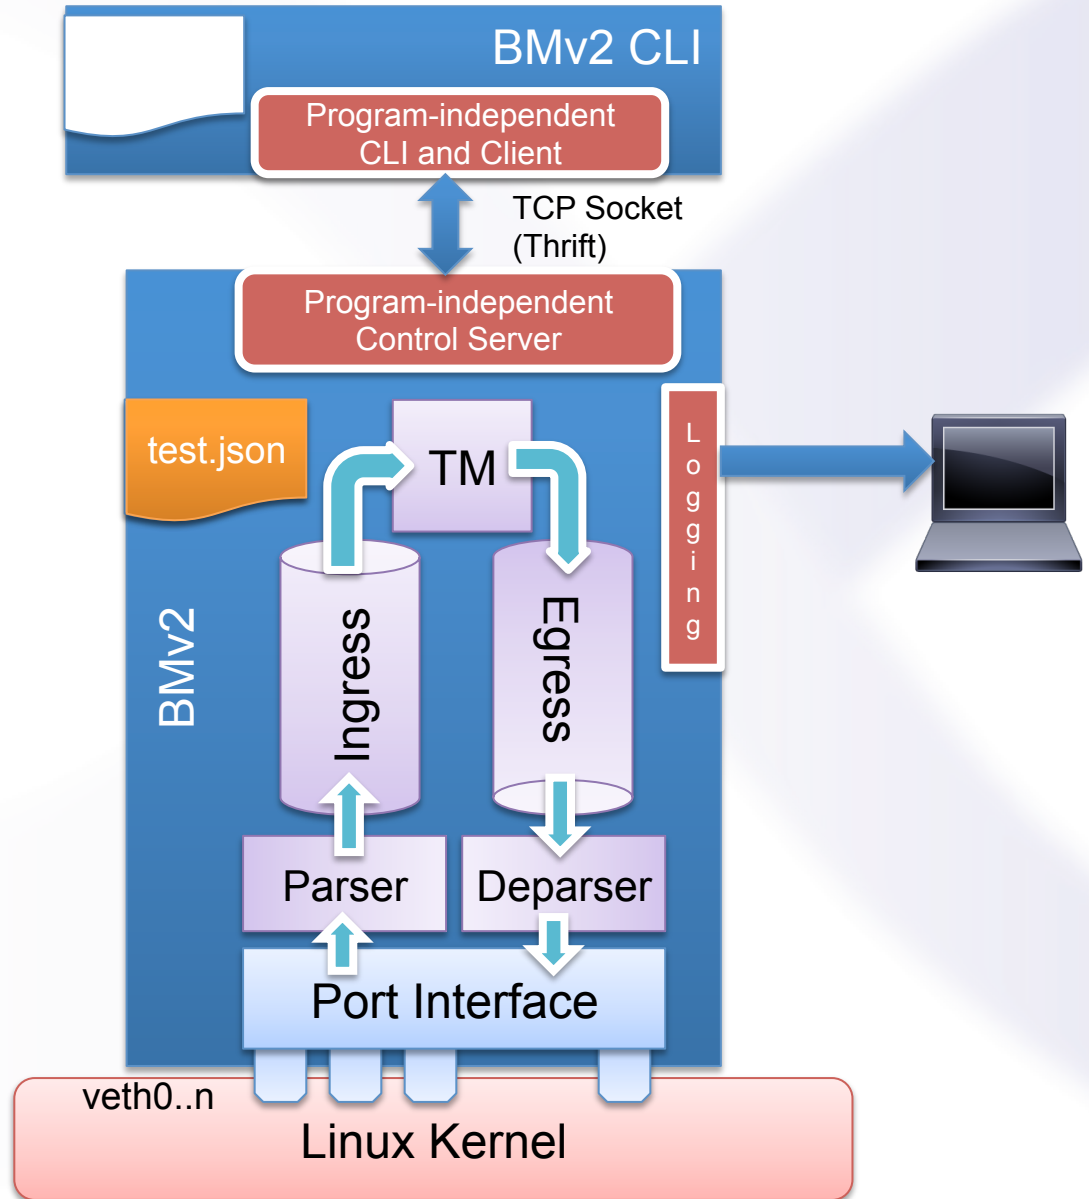

Step 3: Starting the model

```
$ sudo simple_switch test.json --log-console \
-i 0@veth0 -i 1@veth2 ... \
--thrift-port 9090
```



Step 4: Starting the CLI

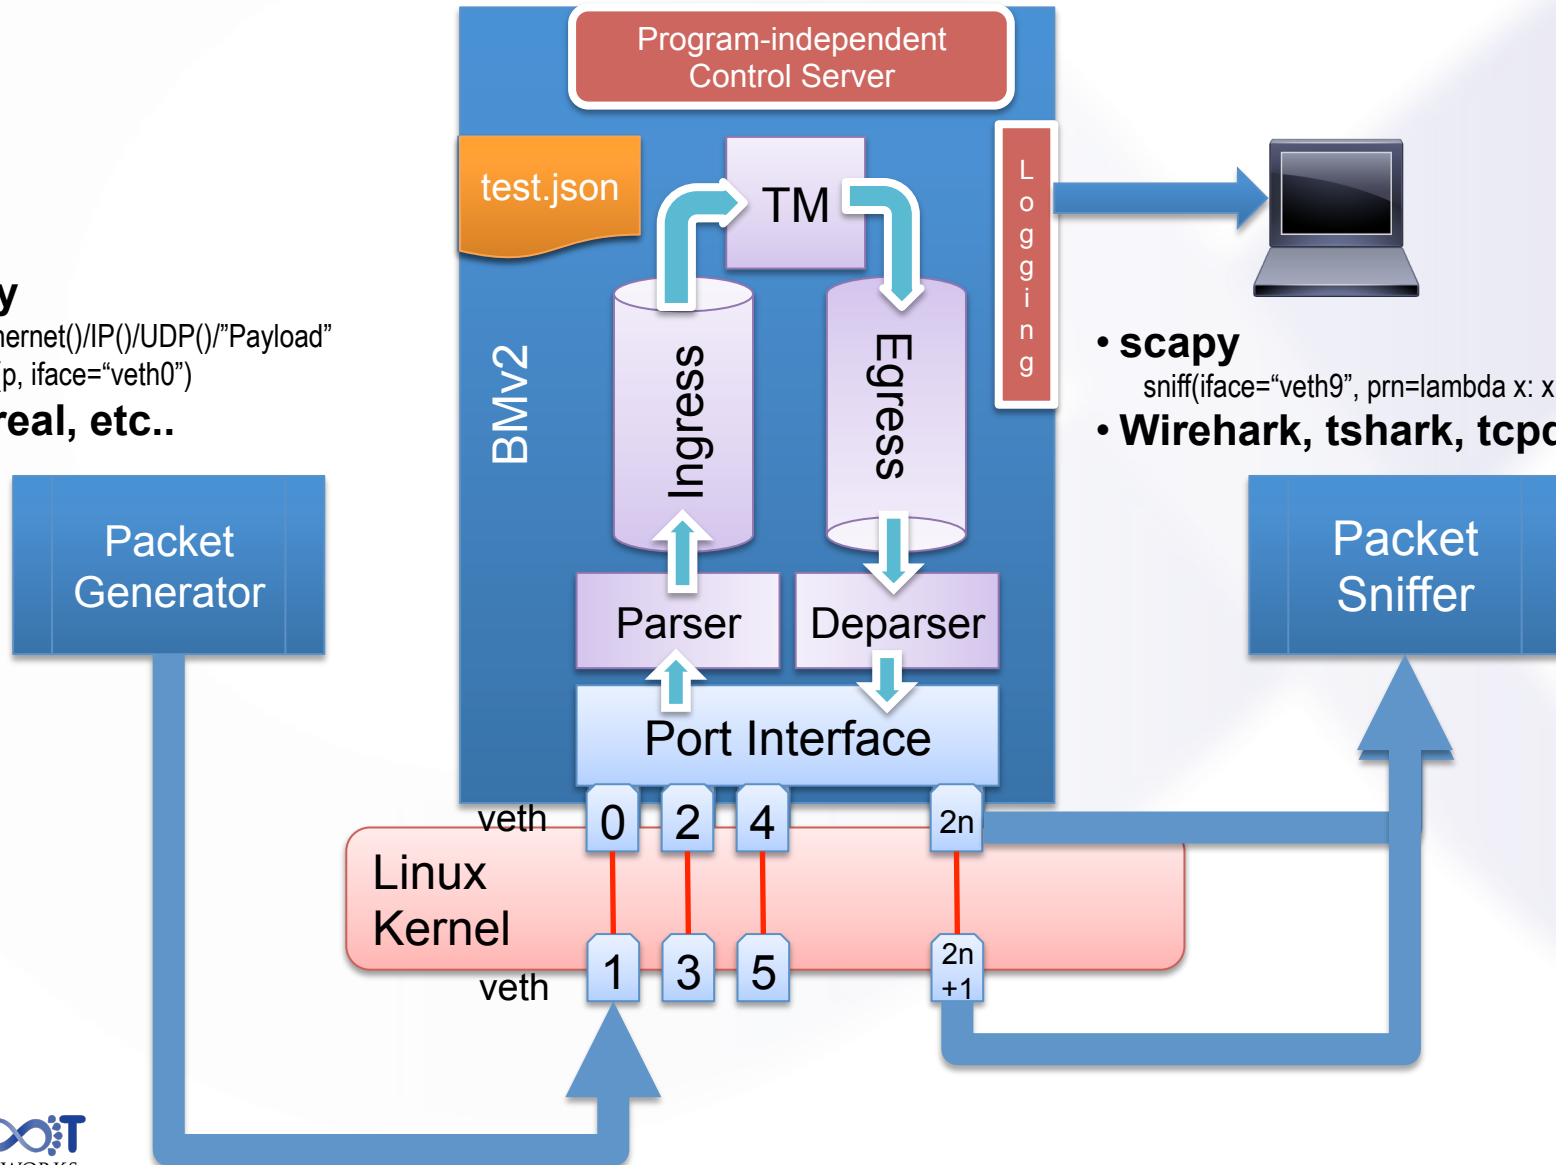
```
$ sswitch_CLI --json test.json
```



Step 5: Sending and Receiving Packets

- **scapy**
`p = Ethernet()/IP()/UDP()/Payload"`
`sendp(p, iface="veth0")`
- **Ethereal, etc..**

- **scapy**
`sniff(iface="veth9", prn=lambda x: x.show())`
- **Wirehark, tshark, tcpdump**

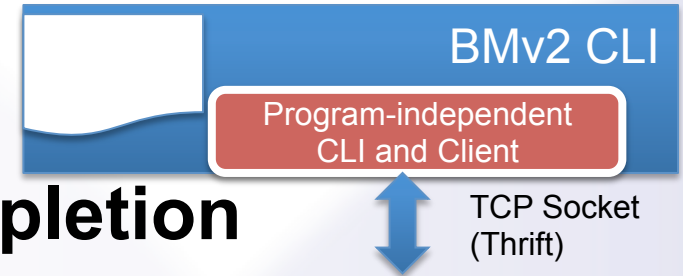


Using the CLI

- Programming the device

Basic Info

- **Simple CLI written in Python**
 - Based on the standard `cmd` module
- **Interactive shell with autocompletion**
- **Simple scripting**
 - Feed a list of commands on STDIN
- **Generic commands for various P4 objects**
 - P4 object definitions are loaded from the JSON file
- **Additional commands for the fixed APIs**
- **No state**
 - Simple translation of commands to Thrift messages
 - Can be restarted (or crashed) without disturbing the model
- **Multiple instances can be started**
 - To communicate with multiple models via separate connections



Getting Help

• Getting the list of commands

RuntimeCmd: **help**

Documented commands (type help <topic>):

=====

counter_read	show_tables
counter_reset	swap_configs
help	table_add
load_new_config_file	table_delete
mc_mgrp_create	table_dump
mc_mgrp_destroy	table_indirect_add
mc_node_associate	table_indirect_add_member_to_group
mc_node_create	table_indirect_add_with_group
mc_node_destroy	table_indirect_create_group
mc_node_dissociate	table_indirect_create_member
mc_node_update	table_indirect_delete
mc_set_lag_membership	table_indirect_delete_group
meter_set_rates	table_indirect_delete_member
mirroring_add	table_indirect_modify_member
mirroring_delete	table_indirect_remove_member_from_group
register_read	table_indirect_set_default
register_write	table_indirect_set_default_with_group
set_queue_depth	table_info
set_queue_rate	table_modify
shell	table_set_default
show_actions	table_show_actions

• Getting the command help

RuntimeCmd: **help table_add**

Add entry to a match table:

table_add <table name> <action name> <match fields> => <action parameters> [priority]

Working with Tables

RuntimeCmd: **show_tables**

```
m_filter          [meta.meter_tag(exact, 32)]
m_table           [ethernet.srcAddr(ternary, 48)]
```

RuntimeCmd: **table_info m_table**

```
m_table           [ethernet.srcAddr(ternary, 48)]
*****
_nop
[]m_action        [meter_idx(32)]
```

RuntimeCmd: **dump_table m_table**

```
m_table:
0: aaaaaaaaaaaa &&& ffffffff => m_action - 0,
SUCCESS
```

Value and mask for
ternary matching.
No spaces around
“&&&”

Entry priority

RuntimeCmd: **table_add m_table m_action 01:00:00:00:00:00&&&01:00:00:00:00:00 => 1 0**

Adding entry to ternary match table m_table

```
match key:      TERNARY-01:00:00:00:00:00 &&& 01:00:00:00:00:00
action:         m_action
runtime data:   00:00:00:05
SUCCESS
```

“=>” separates the
key from the action
data

entry has been added with handle 1

RuntimeCmd: **table_delete 1**

All subsequent
operations use the
entry handle

Packet Replication (Multicast)

Multicast Group (M)

Node 1 (RID 1)

- Port 1₁
- Port 1₂
- ...
- Port 1_p

Node 2 (RID 2)

- Port 2₁
- Port 2₂
- ...
- Port 2_Q

...

Node N (RID N)

- Port N₁
- Port N₂
- ...
- Port N_R

```
RuntimeCmd: mc_mgrp_create 1
```

```
Creating multicast group 1
```

```
SUCCESS
```

```
RuntimeCmd: mc_node_create 10      1 2 3 4 5
```

```
Creating node with rid 10 , port map 111110 and lag map
```

```
SUCCESS
```

```
node was created with handle 1
```

```
RuntimeCmd: mc_node_create 12      10 9 4 6
```

```
Creating node with rid 12 , port map 1011100000000000 and lag map
```

```
SUCCESS
```

```
node was created with handle 2
```

```
RuntimeCmd: mc_node_associate 1 1
```

```
Associating node 1 to multicast group 1
```

```
SUCCESS
```

```
RuntimeCmd: mc_node_associate 1 2
```

```
Associating node 2 to multicast group 1
```

```
SUCCESS
```

A node can also be associated with multiple multicast groups

Managing Mirror Destinations

- **Mirror Destinations (Clone Specs) are used by P4 primitive actions:**
 - `clone_ingress_pkt_to_ingress(clone_spec, field_list)`
 - `clone_ingress_pkt_to_egress(clone_spec, field_list)`
 - `clone_egress_pkt_to_ingress(clone_spec, field_list)`
 - `clone_egress_pkt_to_egress(clone_spec, field_list)`
- **Clone spec is an integer number, representing a “special destination”**

RuntimeCmd: `mirroring_add 12345 2`

- Packets set to clone spec 12345 will go to the switch port #2
- Typical application: designating a certain port for CPU

Scapy – Packet Sniffer and Generator

- **Free Software**
 - <http://www.secdev.org/projects/scapy/>
- **Implemented in Python**
- **Can be imported as a module**
- **Extensible**
 - New packet formats are easily defined
- **Easy to use**
 - Reasonable defaults everywhere
 - Simple Syntax

Simple Examples

- **Creating a packet**

- `p = Ether()/Dot1Q()/IP()/UDP()/("A" * 64)`
- `p = Ether(src="00:00:00:00:00:01", dst="ff:ff:ff:ff:ff:ff") /
Dot1Q(pri=6, vlan=23) /
IP(src="192.168.1.1", dst="192.168.1.255") /
UDP(sport=7, dport=7) / "Vladimir"`

- **Packet display**

- `p.show()`
- `p.show2()`
- `hexdump(p)`

- **Sending the packet**

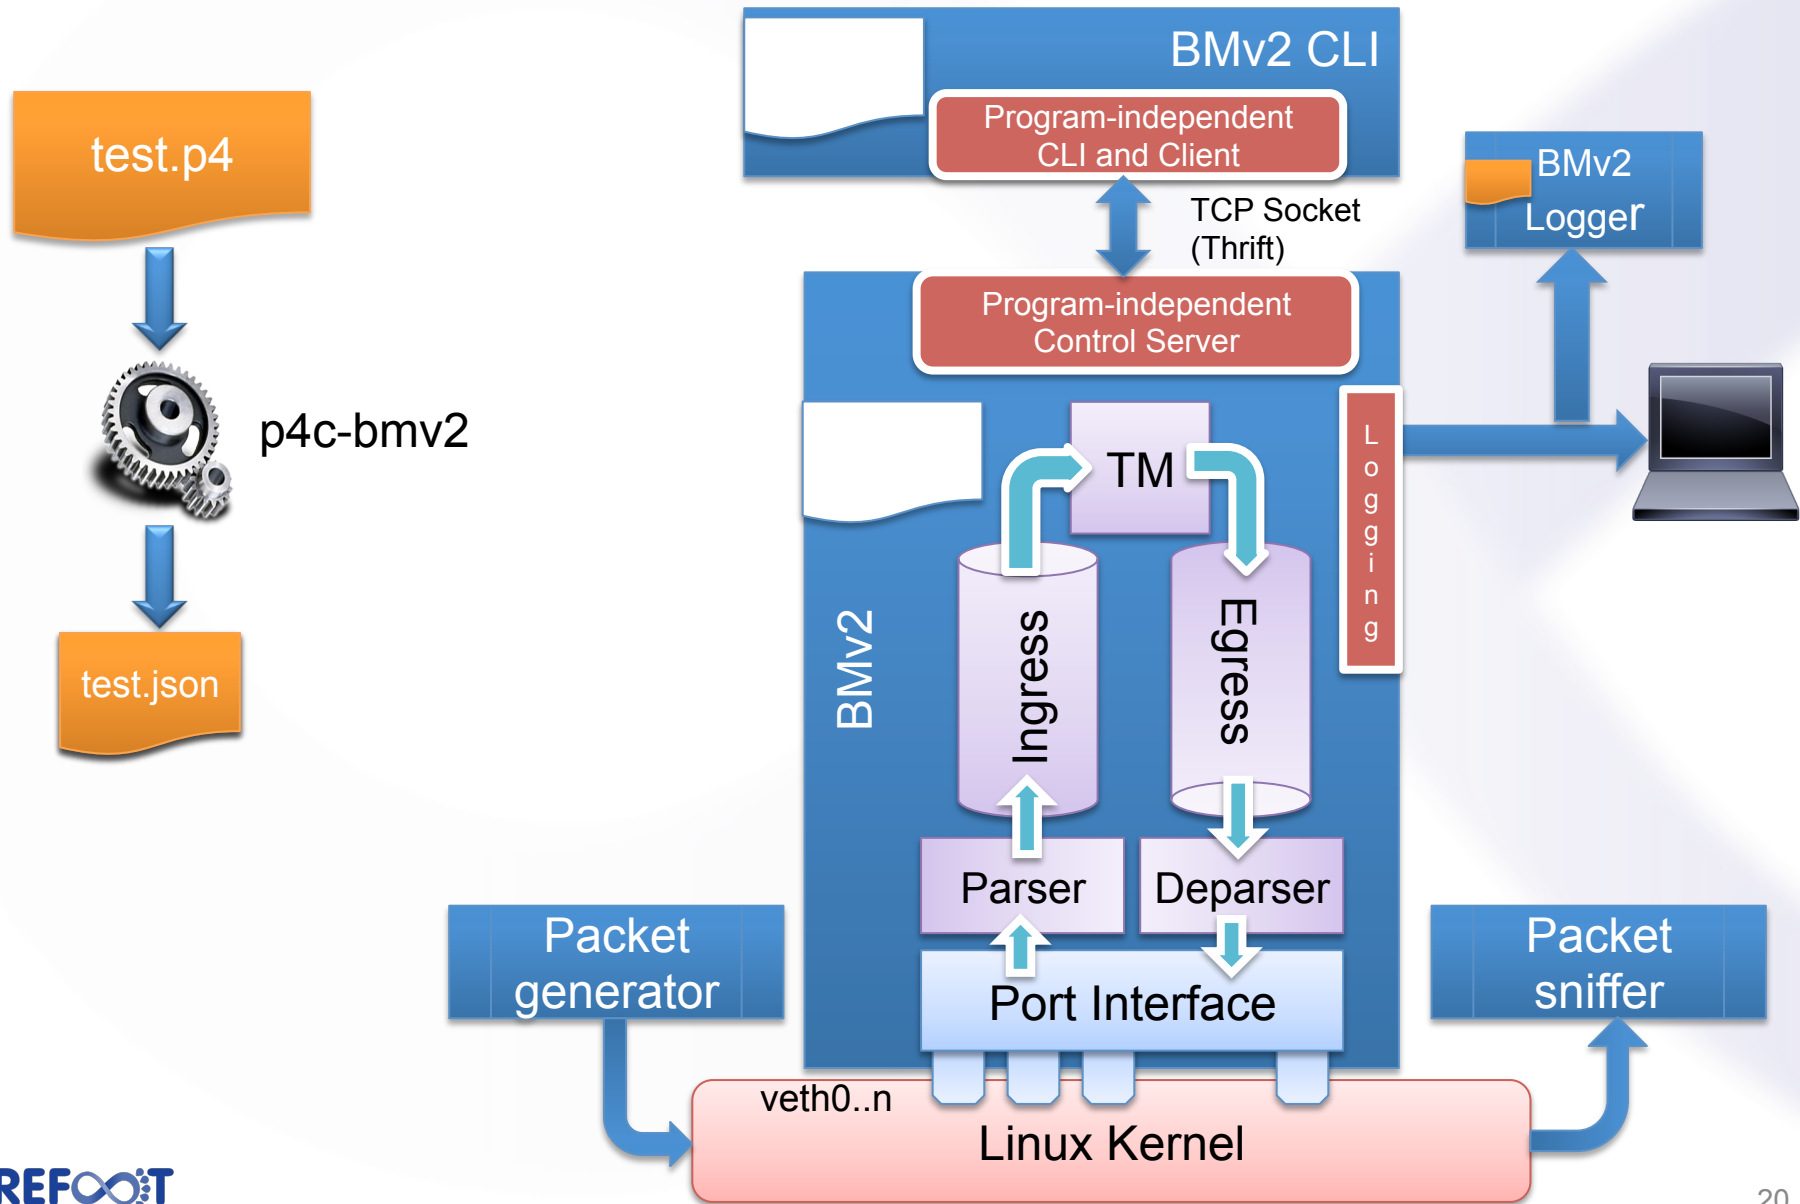
- `sendp(p, iface="eth0", [count=100])`

- **Sniffing**

- `sniff(iface="eth0", prn=hexdump)`
- `sniff(iface="eth0", prn=lambda p: p.show())`

Thank you 8

Basic Workflow



Basic Workflow. P4 Compilation

