

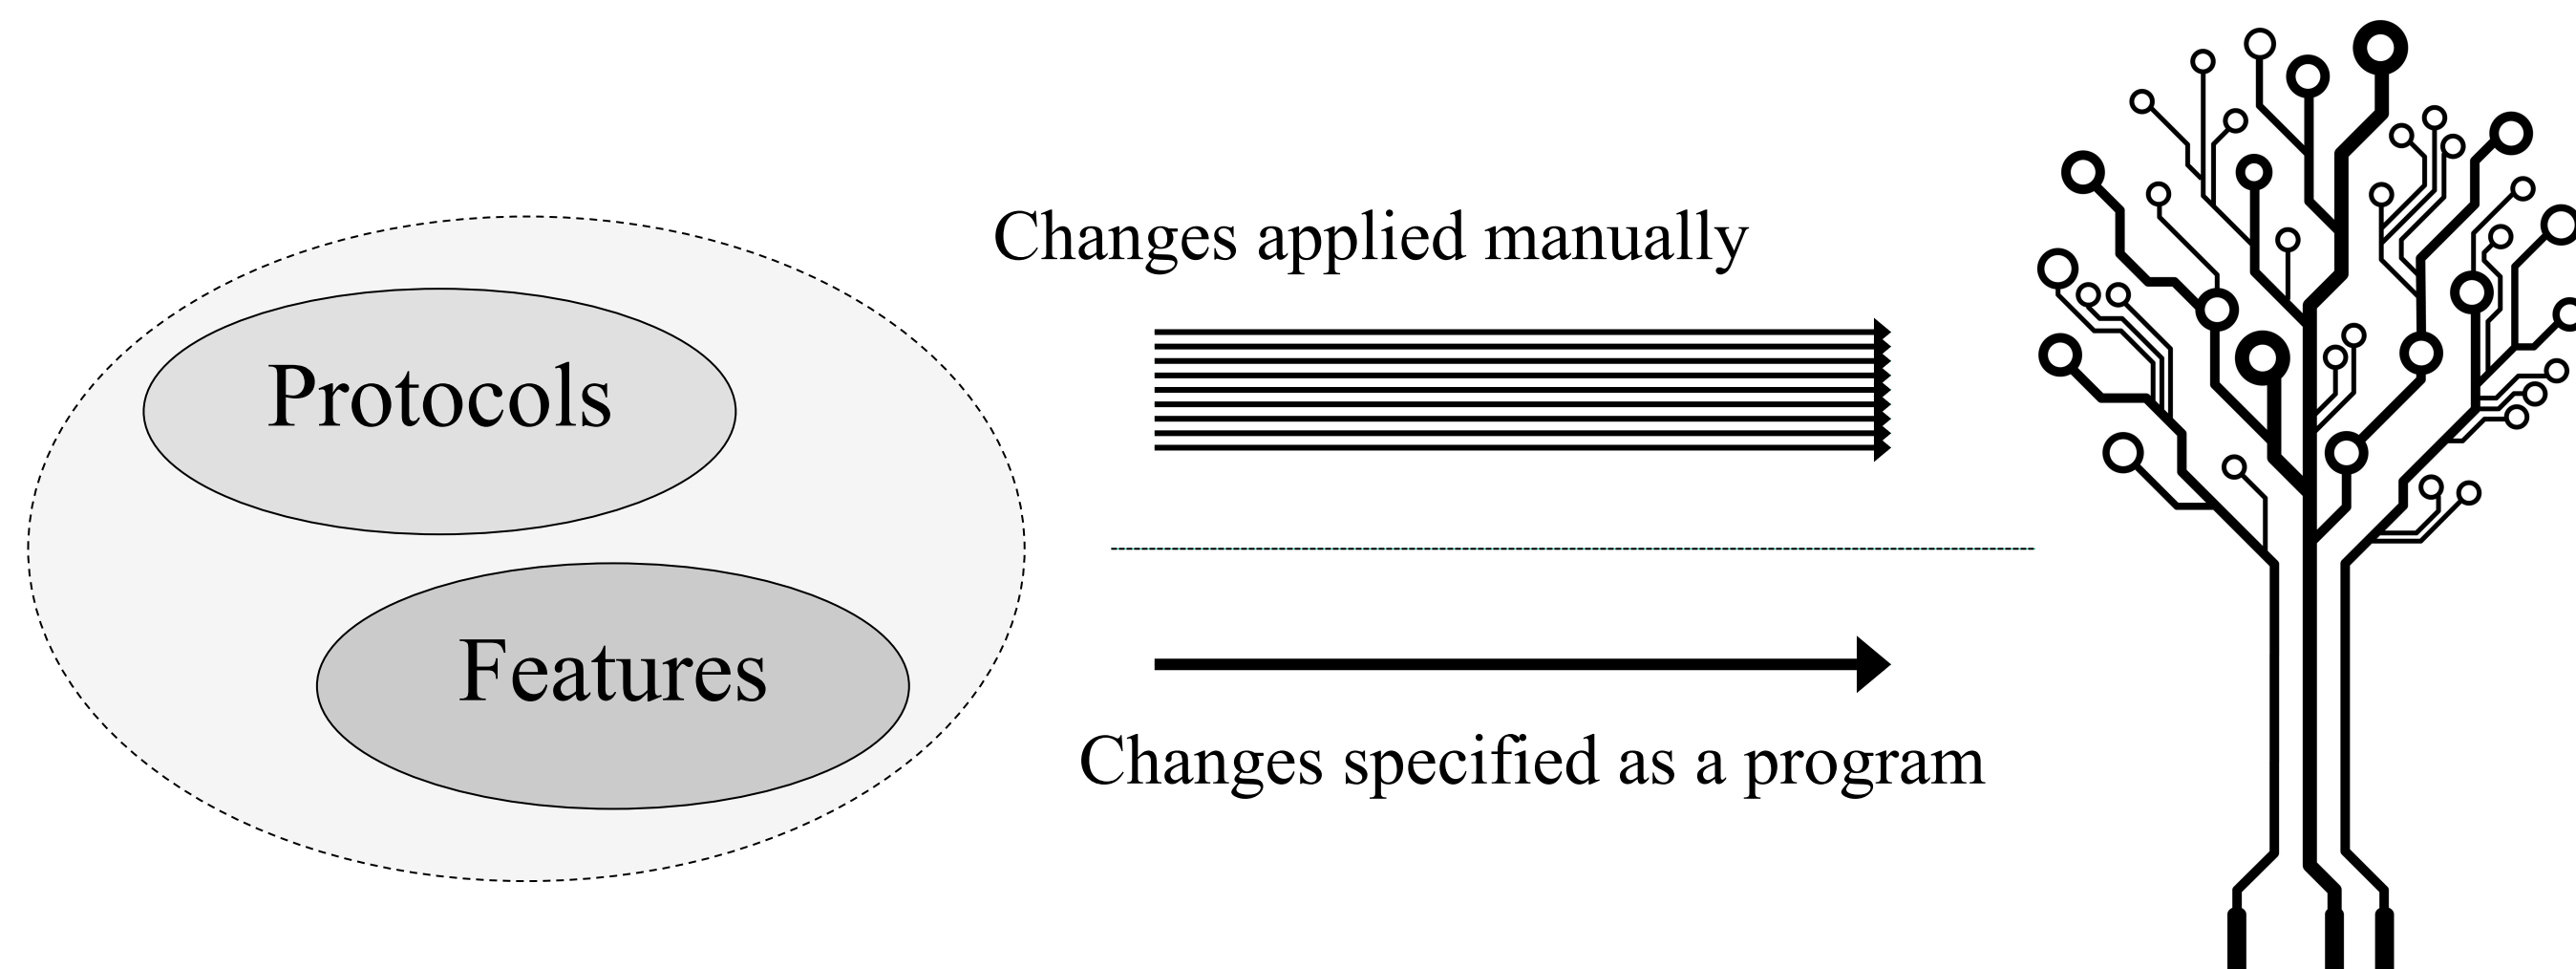


Customizing Open vSwitch using P4

Muhammad Shahbaz, Sean Choi, Ben Pfaff, Chaitanya Kodeboyina, Changhoon Kim,
Nick McKeown, Nick Feamster, and Jennifer Rexford

1. Problem Statement

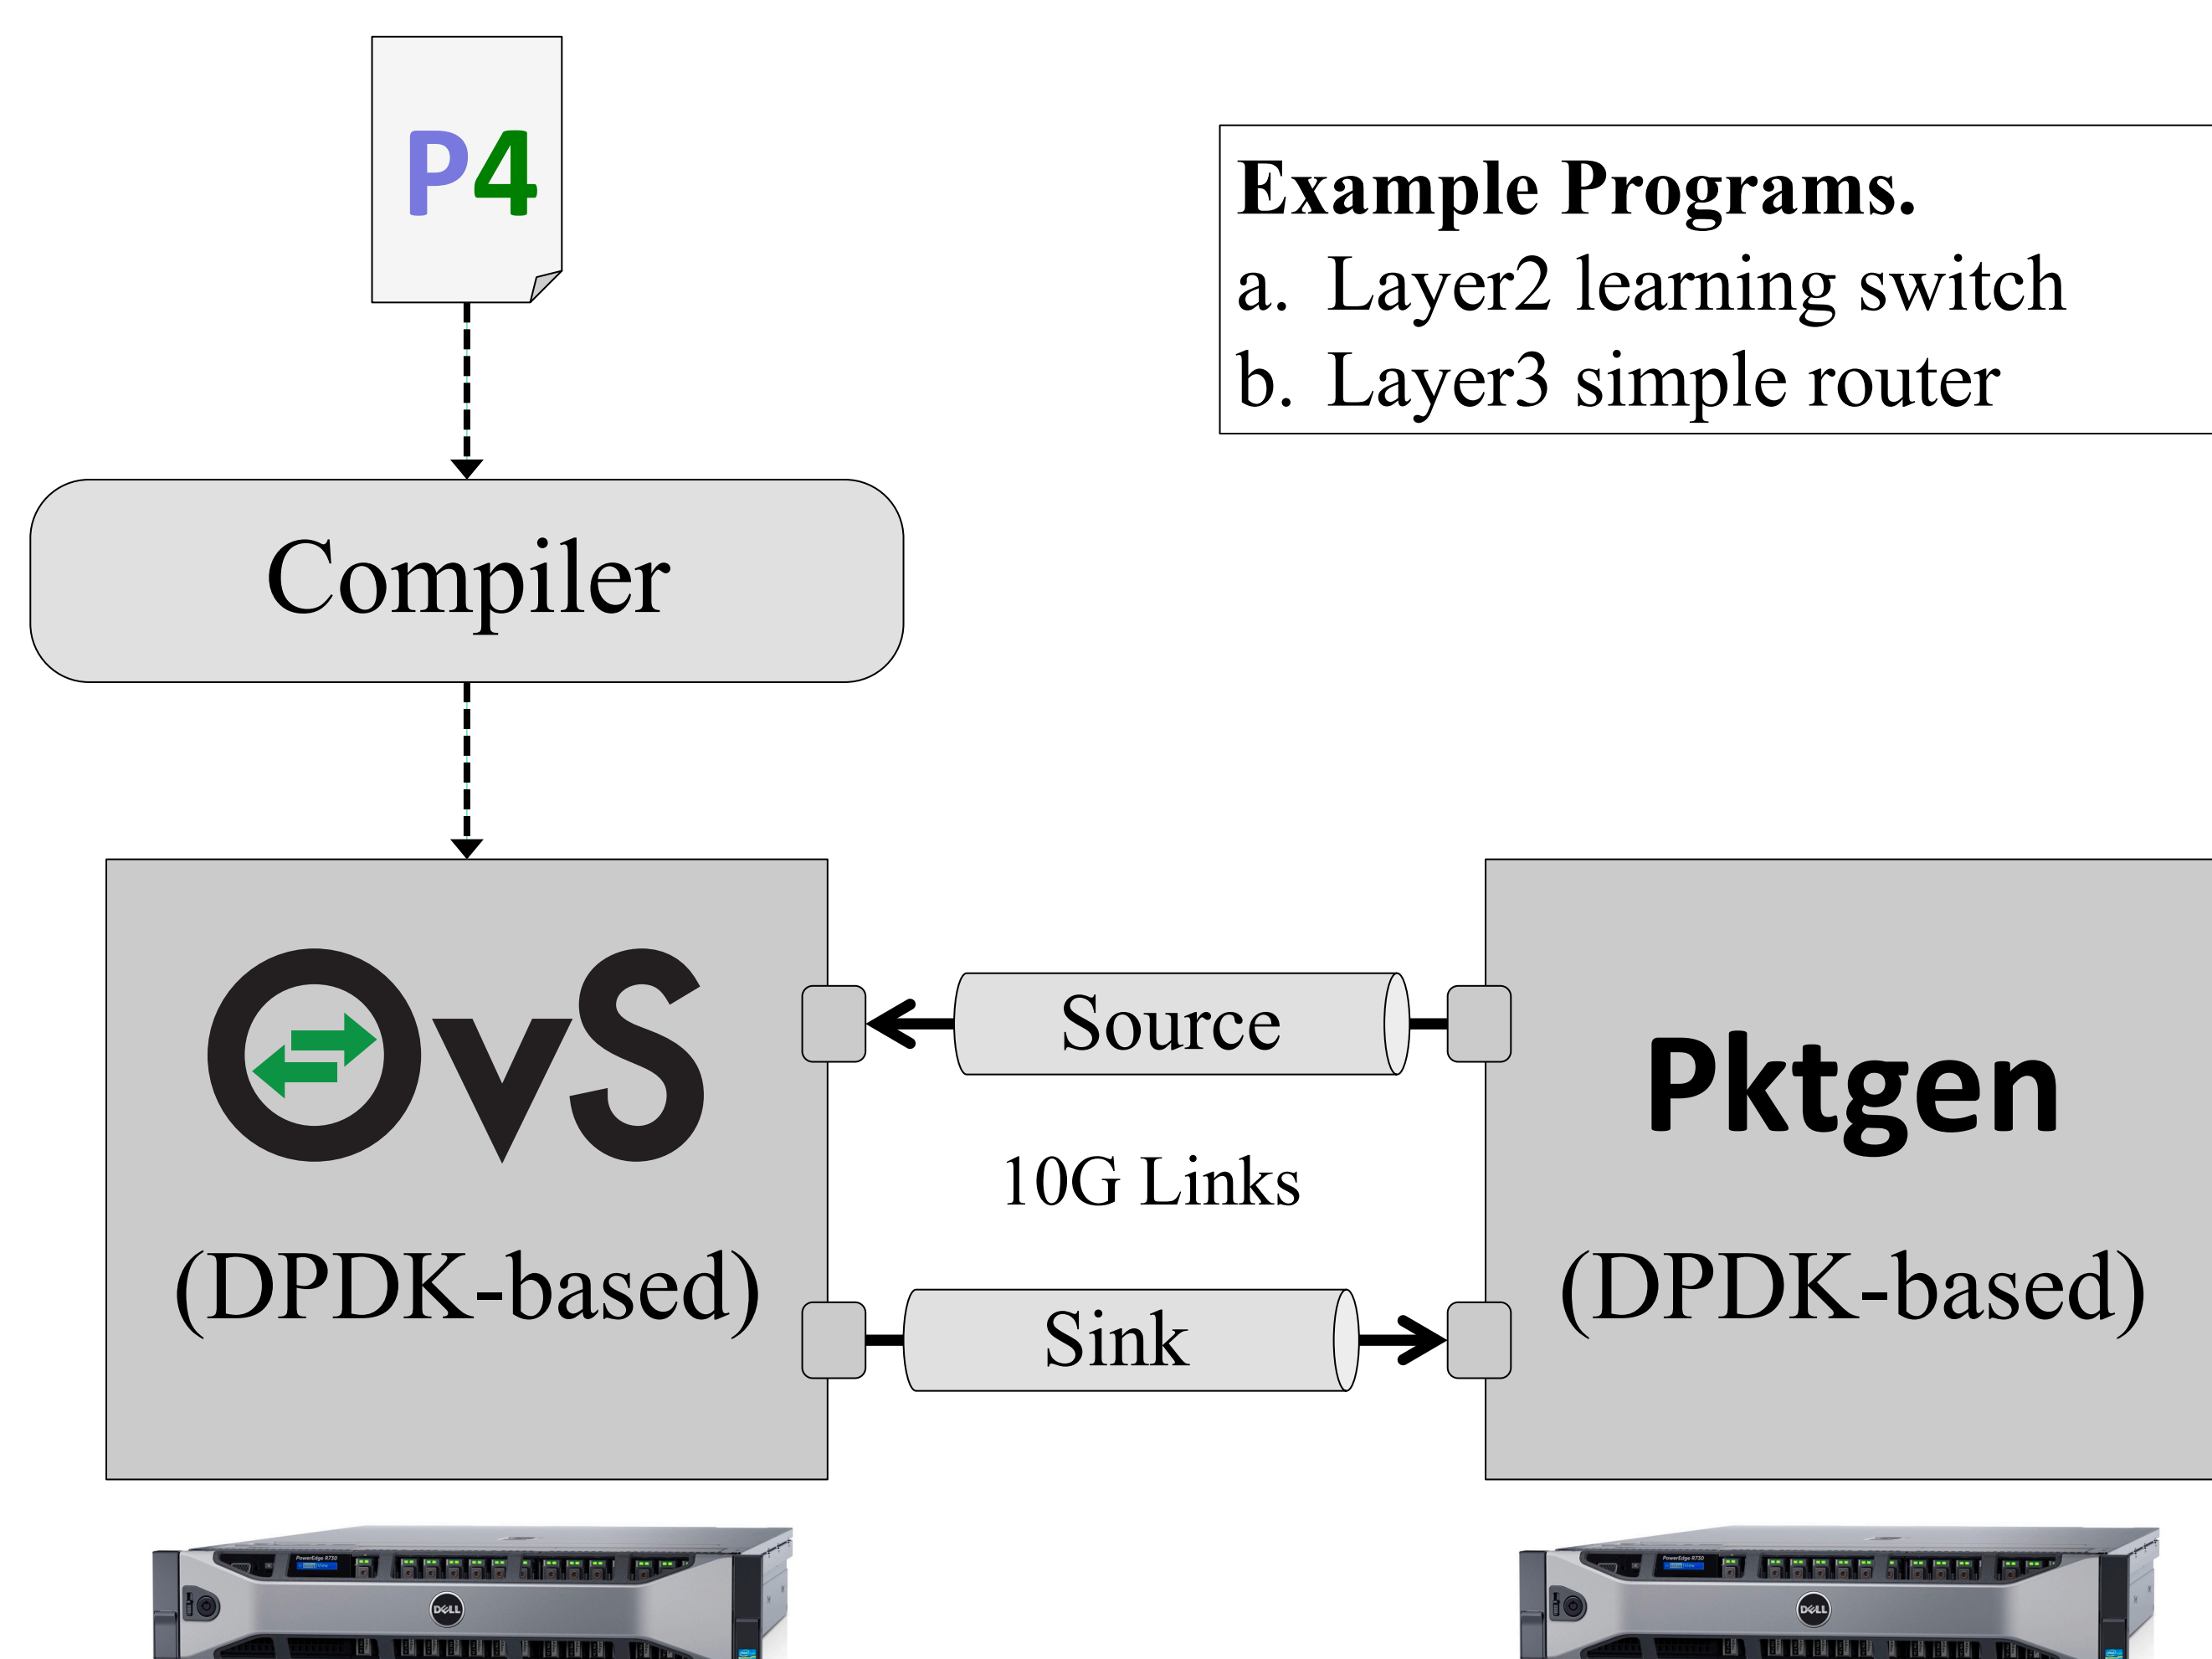
- Adding protocols and features requires
 - manual changes throughout the source tree
 - maintaining changes across newer versions



- Instead, describe these changes as **P4** programs. This has many benefits:
 - Rapid addition of new protocols and features
 - Automated testing and debugging
 - Backward compatibility

What's the Cost of Programmability on Performance?

3. Demo Setup



Example Programs.
a. Layer2 learning switch
b. Layer3 simple router

Server Specifications.

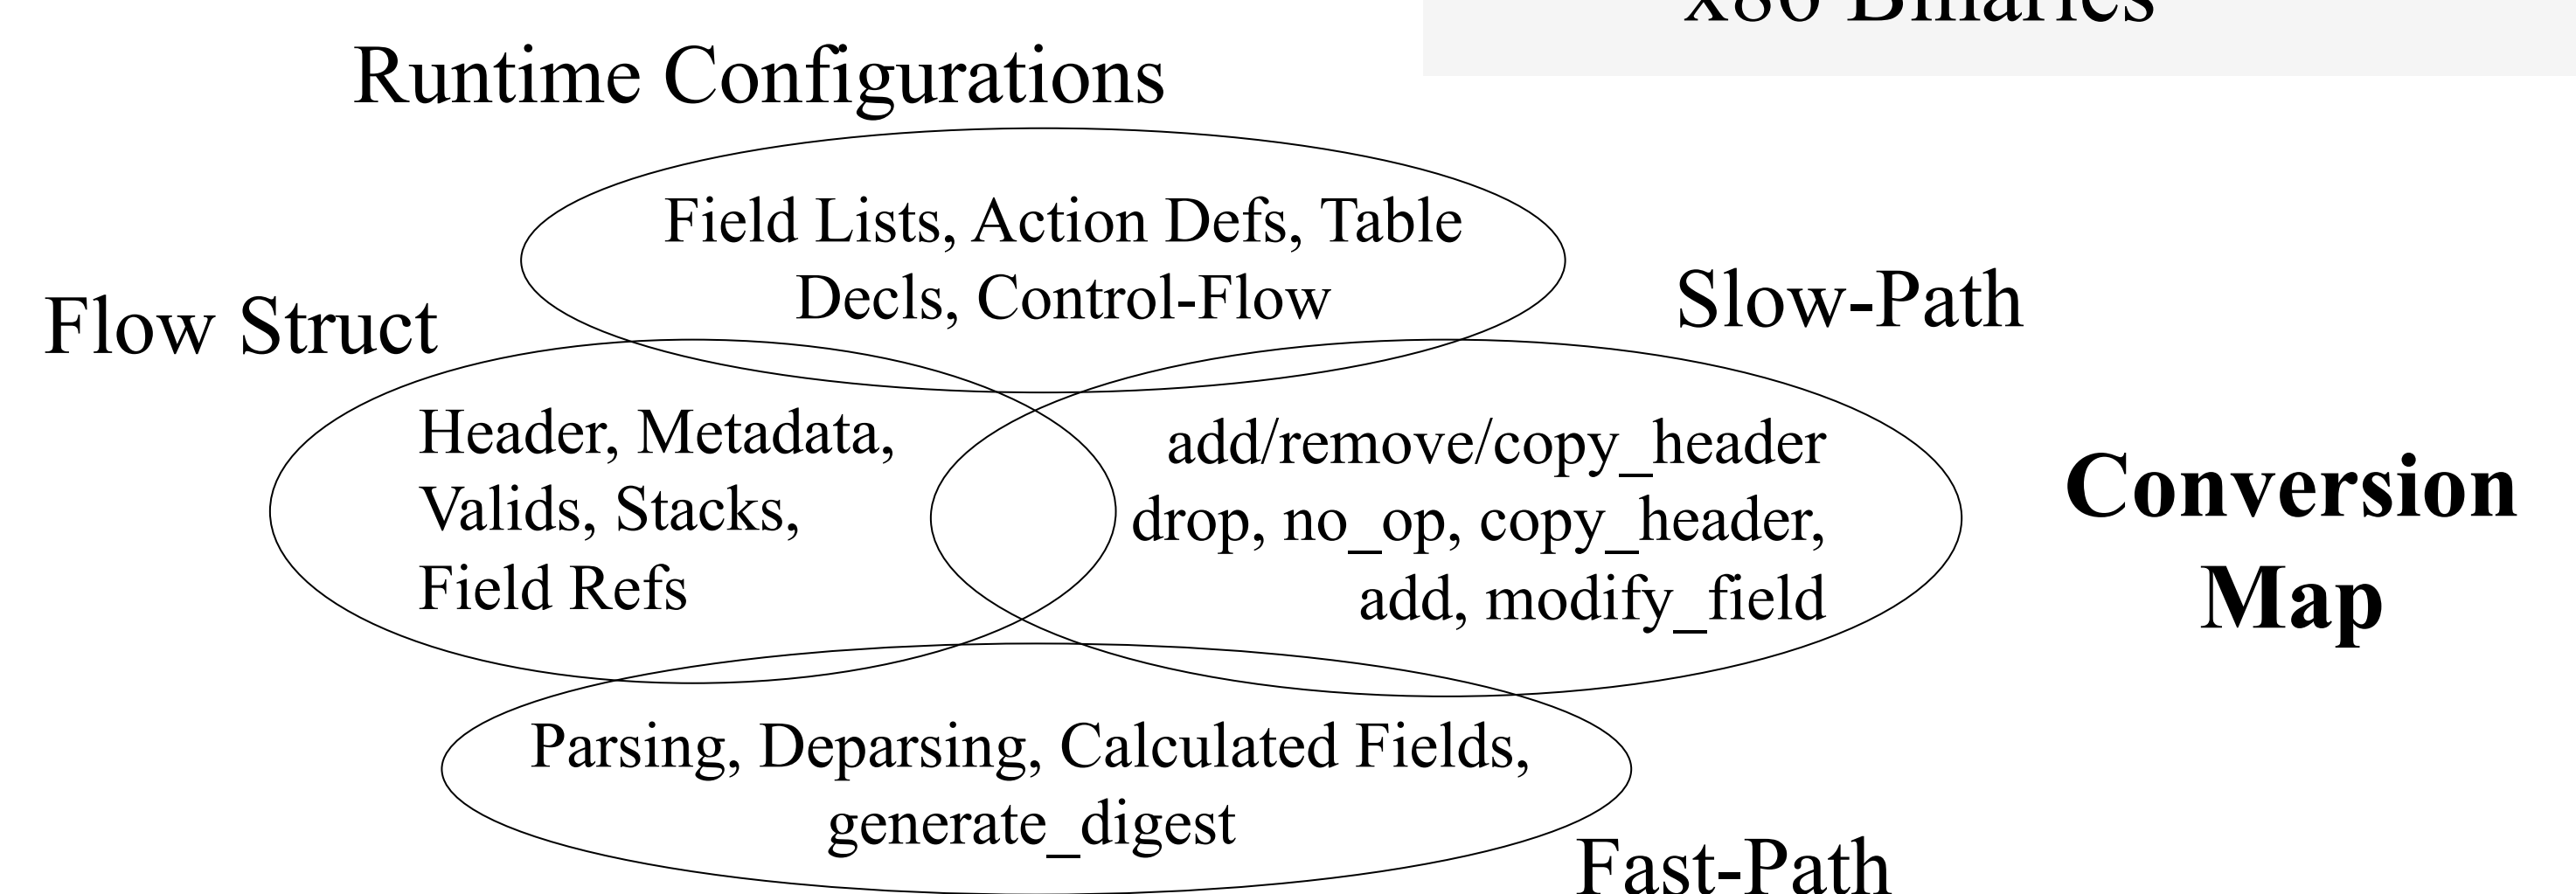
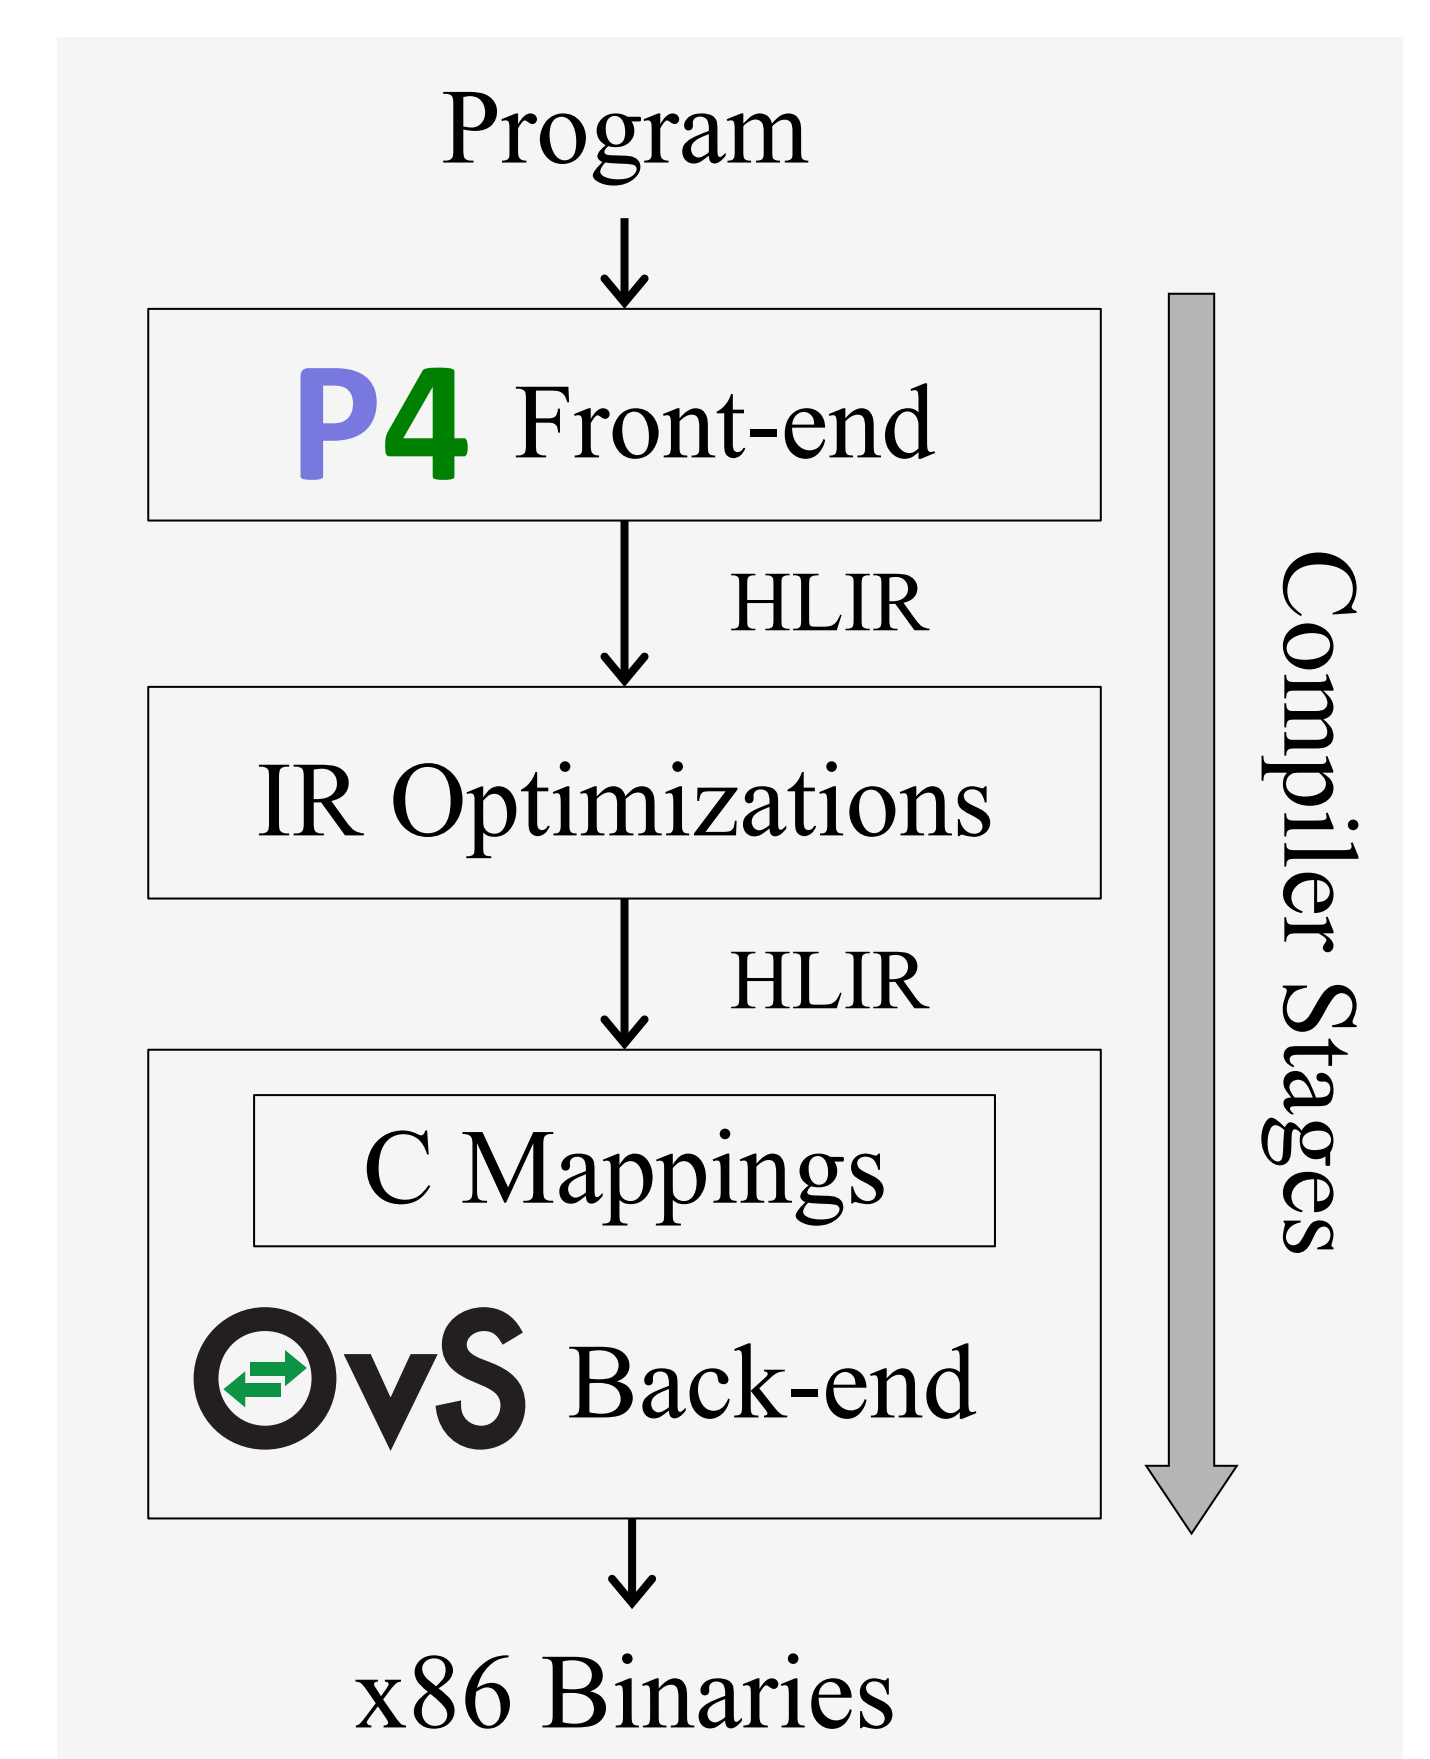
CPU: Intel Xeon E5-2640 v3 2.6GHz
Memory: 32GB RDIMM, 2133 MT/s, Dual Rank
Harddisk: 1TB 7.2K RPM NLSAS 6Gbps
NICs: Intel X710 DP/QP DA SFP+ Cards

2. Approach

- Build a compiler for P4 to Open vSwitch (OvS).

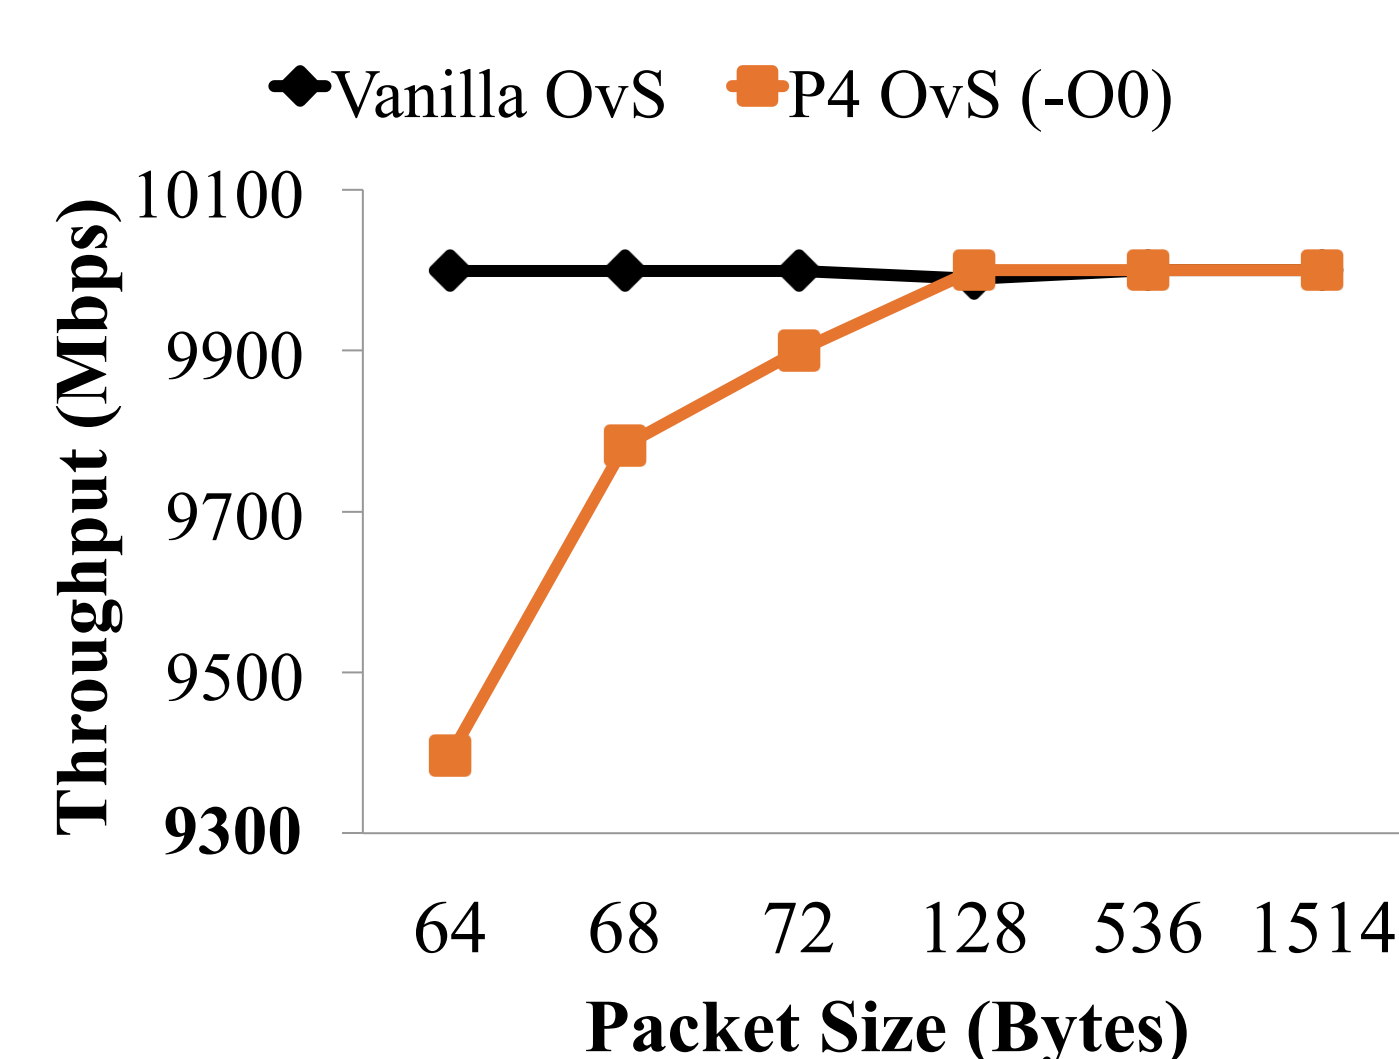
Challenges:

- Mapping P4 to OvS
 - Headers, Parser, and Action Primitives etc.
- Efficient Compilation
 - Program Analysis
 - Optimizations

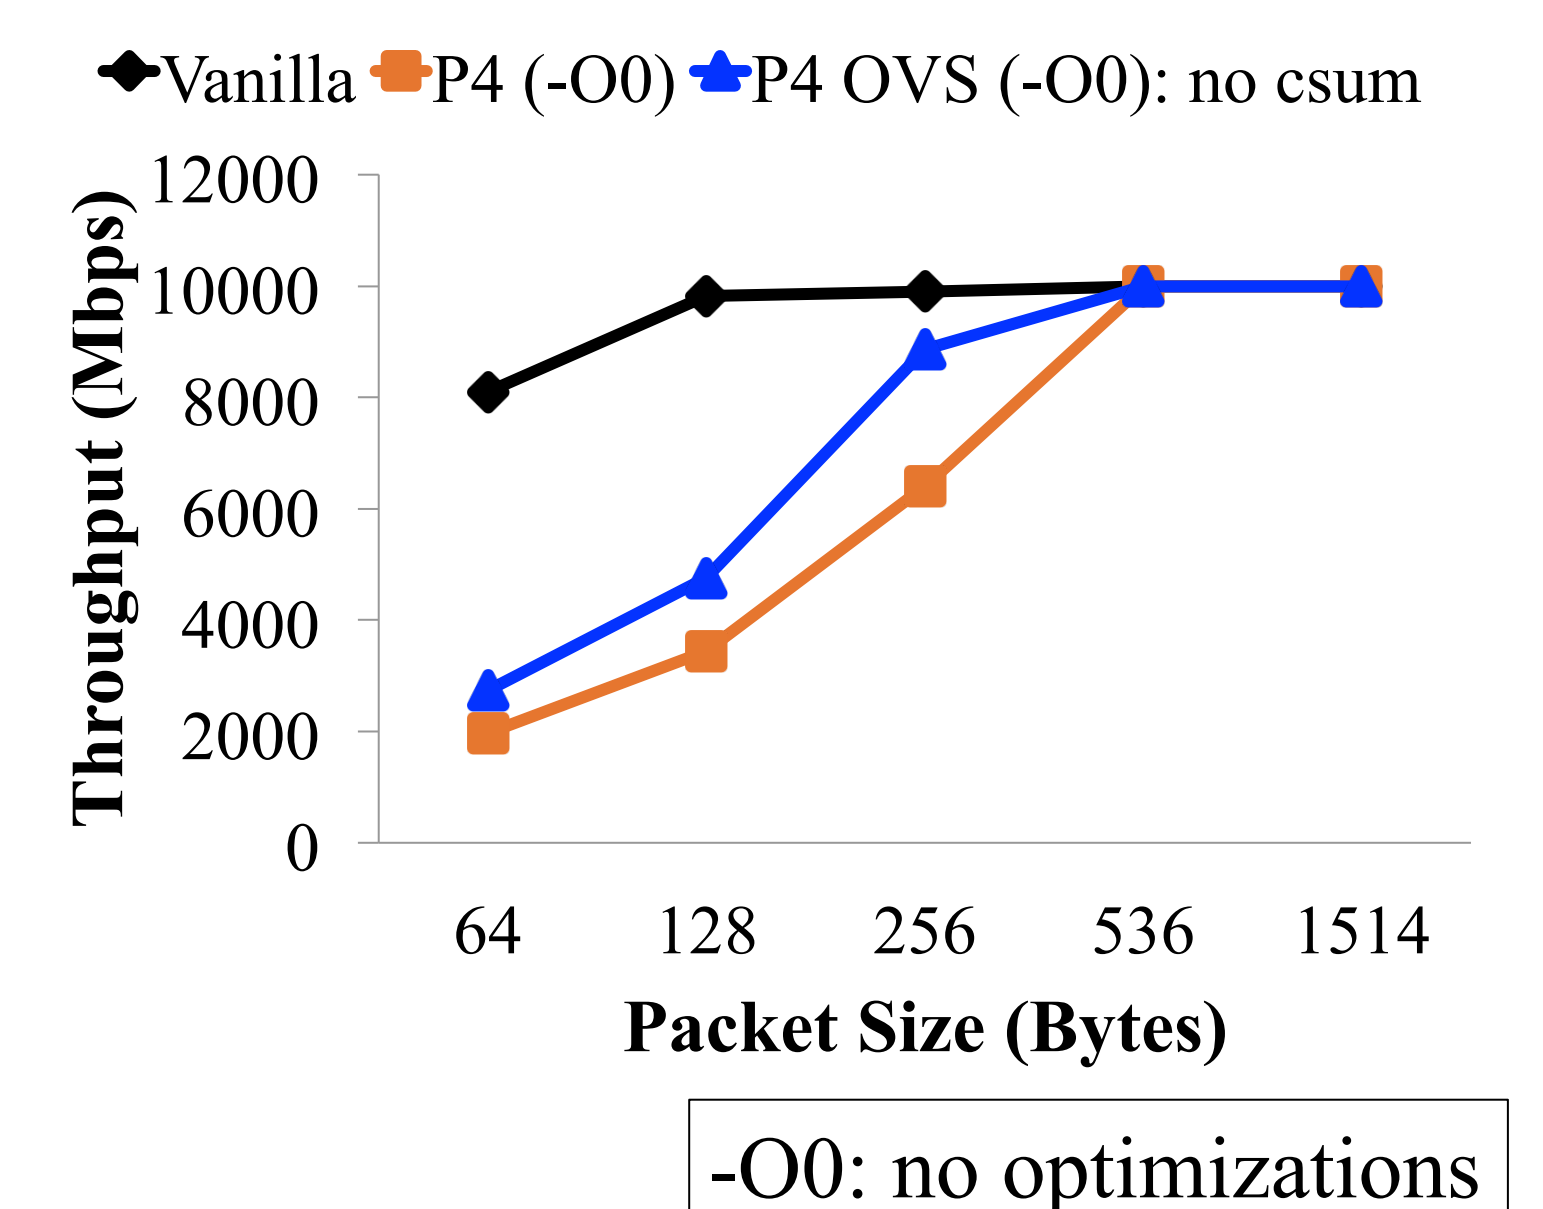


4. Performance Results




a. Layer2 learning switch



b. Layer3 simple router



- Packet parsing is the primary bottleneck
 - Performance degrades with increasing no. of protocols
 - e.g., 19% decrease in Vanilla OvS with IP protocol parsing
 - More significant in P4 OvS as fast-path maintains its own copy of the parsed representation (i.e., extra overhead on parsing)
- Future Goals:**
 - Implement optimizations to avoid extra copies of the parsed representation in the fast-path (based on dead-code elimination and liveness analysis etc.)



Acknowledgments:

The authors wish to thank Mihai Buidu, Ramkumar Krishnamoorthy, and Ed Doe for helpful early discussions and valuable feedback.